About Donna

WHAT IS DONNA?

Donna is a friendly, focused frontend built on top of **LQ-AI**, an open-source legal operating system. Where LQ-AI exposes the full power of a legal AI platform, Donna presents a streamlined experience designed for day-to-day legal work — clean, approachable, and tuned for the tasks that matter most.

The **Powered by LQ-AI** callout at the top of every page in this guide links to a deeper explanation of the backend that drives Donna's answers, document intelligence, and knowledge retrieval.

A TOUR OF THE SIDEBAR

Everything in Donna is reachable from the left sidebar. Here is what each entry does:

- **Assistant** — the main chat interface; start a conversation, attach files, and apply skills or prompts from your library.
- **Projects** — organise your work by matter; chats and documents can be scoped to a project so context stays together.
- **Workflows** — skills, playbooks, saved prompts, and automations that run or accelerate recurring tasks.
- **Tabular** — structured document review; extract and compare fields across a set of files in a grid.
- **Research** — search U.S. case law from CourtListener: find decisions by topic or name, read full opinions, and verify citations.
- **About** — this guide.
- **Settings** — account, model routing, and provider configuration.

HOW THIS GUIDE IS ORGANISED

The rail on the left of this guide has one page per area of Donna. Start with the **Assistant** page — it covers the core chat loop that underlies everything else. From there, work through **Projects**, **Workflows**, **Automations**, and **Tabular** in any order that suits your workflow. The **Research** page covers case-law lookups, **Tools & connections** covers external tools and the in-chat approval flow, and the **Knowledge**, **Models**, and **Trust & citations** pages cover the infrastructure behind Donna's answers.

A STANDING REMINDER

Donna's answers are **not legal advice**. They are AI-generated responses that can contain errors or omissions. Always apply your own professional judgment, and verify any information before relying on it in a legal matter.

The Assistant

STARTING A CONVERSATION

The **Assistant** is Donna's core interface. When you open the app you land on a central composer — a single input field that greets you by name. Type your question or task there and press **Enter** (or the send arrow) to start a new chat. Donna streams the answer back in real time and the conversation moves into its own chat page, where you can continue sending follow-up messages using the composer at the bottom of the screen.

Every chat lives at its own URL, so you can bookmark it, share the link internally, or return to it later from the sidebar's chat history. Within a chat the full message thread is preserved and the composer stays ready at the bottom for follow-up turns.

✦ ENHANCE — REWRITING A VAGUE DRAFT

The **Enhance** button (the sparkle icon labelled **Enhance** in the composer toolbar) takes what you have typed so far and rewrites it into a fuller, more precise prompt before you send it. It is particularly useful when you have a rough idea but are not sure how to phrase it for best results.

The flow is three steps:

- **Run** — click **Enhance**. The button briefly shows "Enhancing…" with a cancel option while the backend rewrites your draft.
- **Preview** — a preview card appears above the text box showing the expanded prompt and, optionally, a collapsible "Why these changes" list explaining each edit.
- **Accept or discard** — click **Use this** to replace the text in the composer with the enhanced version, or **Discard** to go back to your original draft. Neither action sends the message; you can still edit before hitting send.

If the backend finds nothing to improve it shows a brief "No changes suggested" note and the composer is unchanged. Enhance is context-aware: if you have skills attached to the message, those skill names are sent along so the rewrite can take your chosen workflow into account.

ATTACHING SKILLS

Skills are structured workflows that guide Donna's response — for example, a contract review checklist or a clause extraction routine. To attach one, click the **Skill** button (a + icon) in the composer toolbar. A small search panel opens; type a keyword to find skills by name or description and click the skill you want to add.

Attached skills appear as pill badges inside the composer. Some skills expose required or optional input fields — those appear inline between the pills and the text area so you can fill them in before sending. Required fields must be filled before the send button becomes active. Remove a skill at any time by clicking the × on its pill. You can attach more than one skill to a single message; each is sent alongside your prompt.

KEEP SKILLS ON

Once you have a skill attached to a message, the **Keep skills on** toggle in the composer toolbar lets you carry that skill — and any others attached — into every follow-up turn of the same chat without re-attaching. When active, a quiet **"Keeping N on"** indicator appears alongside the toggle; hover it to see the skill names that are being carried forward.

The toggle is **off by default** and scoped to the current chat — opening a different chat reflects that chat's own setting, and a brand-new chat always starts off. Turning it off on any turn clears the sticky set immediately. Every turn still records exactly which skills it used in the receipts drawer, so the audit trail is always honest even when skills are carried automatically.

ATTACHING FILES

To ground an answer in a specific document, click the **paperclip** button in the toolbar or drag a file directly onto the composer. You can attach up to **16 files per message**. Each file appears as a pill showing the filename and its upload status ("uploading" → "pending" → "processing" → "ready", or "failed"). The send button stays disabled until every attached file has finished processing — a failed file must be removed before you can send. Remove any file at any time with the × on its pill.

Files attached to a message are sent to the backend alongside your text and are available to the model for that turn. If you need to make a document available across many chats, upload it to a project's knowledge base instead — see the **Knowledge** page of this guide for details.

SAVED PROMPTS

The **Prompts** button (bookmark icon) opens your saved prompt library. Click a saved prompt to insert its text at the cursor position in the composer — it splices in without overwriting what you have already typed. You can also save the current draft directly from the panel: click **+ Save current draft as a prompt**, give it a name and optional tags, and save it for future reuse.

CHOOSING A MODEL

The model picker in the composer toolbar shows the currently selected alias (for example **smart**) alongside the resolved model it maps to (for example "Opus 4.7"). Click it to open a menu grouped into **Cloud** and **Local** sections. Choose any model for the next message; your selection persists across turns in the same browser session.

When a chat is scoped to a project that requires a minimum capability tier, lower-tier models appear greyed out in the picker and cannot be selected. Donna automatically promotes your chosen model to the floor tier in that case so the answer meets the project's quality requirements. The **Models** page of this guide explains tiers, cloud vs. local routing, and how to configure model preferences in Settings.

READING THE ANSWER: CITATIONS

When Donna's answer draws on documents in a knowledge base or on files you attached, it embeds numbered citation pills inline in the text — small superscript markers like [1] or [2]. Hover a pill to see a popover with the source text excerpt, the document filename, and the page number if available. The popover header is colour-coded: green for verified citations, amber for those with caveats, and red for citations that could not be matched to a source.

Click a citation pill (or press Enter/Space when it is focused) to open the **Document panel** on the right side of the chat, which shows the full source document with the cited passage highlighted. This lets you read the answer and its source side by side without leaving the chat. The **Trust & citations** page of this guide explains how citations are verified and what each colour state means.

RECEIPTS DRAWER

Every chat exposes a **Receipts** button in the top-right corner of the chat header. Click it to open the receipts drawer, which lists every backend event that occurred during the conversation — retrieval steps, skill invocations, model calls, and more — in chronological order. Use the kind-filter chips at the top of the drawer to show or hide event categories. You can also export the full event log as a JSONL file via the **Export** link in the drawer header, useful for debugging or auditing a response.

Projects

WHAT IS A PROJECT?

A **Project** groups related chats, files, knowledge bases, skills, and standing context under one name. Instead of starting every conversation from scratch, you give Donna a shared foundation — background notes, uploaded documents, and relevant knowledge — that applies automatically to every chat inside that Project.

A note on naming: the sidebar entry is labelled **Projects**. In some parts of the interface (page headings, form labels, archive dialogs) you will see the word **matter** instead. These refer to the same thing; the terms are interchangeable.

CREATING A PROJECT

Open **Projects** in the sidebar. Click **New matter** in the top-right corner. A dialog asks for:

- **Matter name** — required. Choose something that reflects the engagement or topic.
- **Description** — optional free text shown beneath the name in the list.
- **Privileged matter** — see below.
- **Minimum model tier** — see below.

Once created, the Project appears in the list sorted by last activity. You can **Rename** it at any time from the Project detail page; you can also **Archive** it to remove it from your active list without deleting any chats inside it.

PRIVILEGE AND MINIMUM MODEL TIER

Checking **Privileged matter** does two things:

- Every chat started inside the Project is flagged as privileged in the audit log.
- A **Minimum model tier** (1–5) becomes required — you must choose one before you can save a privileged Project. Higher tiers require cloud-hosted models.

Privileged Projects display a **Privileged** badge — a small lock-icon pill — next to the Project name in the list and on the Project detail page. Non-privileged Projects can optionally set a minimum tier without enabling the privileged flag.

ATTACHING FILES

The **Files** section on a Project's detail page lets you attach documents directly to the Project. When no files have been added yet, a drop zone is shown — drag files onto it or click to open the file picker. After the first upload an **+ Add file** button appears below the file list so you can keep adding more. Uploaded files are available to chats in that Project and to any linked knowledge base workflows.

KNOWLEDGE BASES AND SKILLS

A Project can be linked to one or more **Knowledge** bases. Linking a knowledge base makes its indexed documents retrievable during every chat in the Project. Use the **Knowledge** section on the detail page to link or unlink bases; the **Manage** link takes you to a knowledge base's own settings.

The **Skills** section lets you attach named skills to a Project. Attached skills appear as removable tags and are available to Donna when answering chats scoped to that Project.

STANDING CONTEXT

The **Context** section holds free-form Markdown that Donna receives at the start of every chat in the Project — no matter who starts it or when. Use it for background facts, formatting preferences, recurring instructions, or any information that would otherwise need to be repeated in every prompt. The field supports up to 100 KiB of text; a live byte counter is shown as you type.

SCOPING A CHAT TO A PROJECT

From the Project detail page, click **+ New chat in this matter**. The chat opens already scoped to the Project: Donna receives the standing context, can draw on linked knowledge and skills, and sees the attached files. The chat list at the bottom of the detail page shows all chats ever started in that Project, with a message count for each.

If a Project is marked privileged, the minimum tier setting is also enforced for those chats, ensuring they are handled by a model at or above the required capability level.

Workflows

The **Workflows** section is the single place in Donna where you save and organise reusable building blocks. One sidebar entry opens a hub with four tabs — **Skills**, **Playbooks**, **Prompts**, and **Automations** — each covering a different kind of reuse.

SKILLS

A skill is a saved block of instructions that shapes how Donna responds to a message. When you attach a skill to the composer, those instructions are automatically included with your next send — you do not need to retype them each time.

The **Skills** page lists your own skills under **Your skills** and a searchable catalogue of built-in skills below. Built-in skills cover common legal tasks and cannot be edited directly, but you can **Fork** any of them to create your own copy — the fork appears in **Your skills** and you can edit it freely.

To author a skill from scratch, press **New skill**. The edit form lets you set a name, description, version, an optional slash command alias (e.g. **/nda**), tags, and the skill body — the instructions Donna receives. You can archive a skill you no longer need; it will be removed from the composer.

Some skills declare **inputs** — structured fields the skill expects you to fill in before sending. When you attach such a skill in the composer, an inputs panel appears inline. Required inputs are flagged; a skill with unfilled required inputs cannot be sent until they are provided. Optional inputs can be expanded separately. Input types include text, integer, boolean toggle, and enum (drop-down).

PLAYBOOKS

A playbook is a set of negotiation positions Donna applies to a contract document. Each position describes a clause topic and the stance Donna should take — preferred language, fallback, and non-negotiable limits — so you get consistent redlines every time instead of rewriting guidance from scratch.

The **Playbooks** page groups available playbooks by contract family. Tap any playbook to open its detail view, which shows the full list of positions in order. From the detail view you can **Apply to a document** (admin accounts only in this version), **Edit** the playbook if you own it, **Duplicate** it to create a private copy, or **Delete** it.

To create a new playbook, press **New playbook** from the list. Two paths are available:

- **Generate from documents** — pick existing matter files and name a contract type; Donna drafts a full playbook from those documents and shows a review editor before you save.

- **Start from scratch** — open the playbook editor directly and build positions by hand.

SAVED PROMPTS

Saved prompts are reusable text snippets accessible from the chat composer. When you find yourself typing the same phrasing repeatedly — a standard instruction, a preferred opener, a recurring question — you can save it once and insert it with two clicks.

The **Prompts** page lists all your saved prompts. Press **New prompt** to create one with a name, body text, and optional tags. You can edit or delete any prompt from the same list.

In the composer, the **Prompts** button opens a small picker above the input area. Type to filter by name or tag, then click a result to insert its text at your cursor. You can also save whatever you have already typed as a new prompt directly from the picker — name it and press **Save** without leaving the chat.

AUTOMATIONS

The **Automations** tab is where Donna runs work on its own — one-off runs, cron-style schedules, and watches that fire when new documents arrive in a knowledge base, each with a full receipt of what the agent did and what it produced. See the dedicated Automations guide.

Automations

Automations let Donna run work on its own — a playbook or a skill executed against a knowledge base in the background, on demand or on a schedule. Every run leaves a full transparency receipt: what the agent did, what it cost, why it stopped, and what it produced.

TURNING AUTOMATIONS ON

Automations are off by default. Enable them with the **Automations** toggle in **Settings → Preferences**. Until then, automation pages show an inline **Automations are off** notice with an enable button, so you can opt in right where you need it.

WHERE TO FIND IT

Automations is the fourth tab of the **Workflows** hub, alongside Skills, Playbooks, and Prompts. Inside it, five views cover the lifecycle: **Sessions** (every run, past and live), **Schedules**, **Watches**, **Notifications**, and **Review**.

STARTING A RUN NOW

Press **Run now** to start a one-off session. Choose what to run — a **playbook** or a **skill** — then pick the **target knowledge base** the run works against (required), optionally link a **matter**, and optionally set a **cost cap** in USD. Donna starts immediately and takes you to the live receipt.

SCHEDULES AND WATCHES

- **Schedules** run the same configuration on a recurring timer. You describe the cadence with a cron expression — the form previews it in plain language before you save.
- **Watches** trigger a run automatically when new documents arrive in a knowledge base — useful for reviewing incoming material without checking manually.

Both are editable after creation — including reassigning or unlinking the matter — and both accept a per-run cost cap.

RECEIPTS: WHAT THE AGENT DID

Open any session to see its receipt. The header shows the run's status, how it was triggered, total cost, and — for finished runs — why it stopped. The **Activity** timeline lists every phase the run moved through and every tool call it made, in order. While a run is live the page updates automatically every few seconds.

RESULTS: WHAT THE RUN PRODUCED

The **Results** section shows the run's work product: its findings, in the order the run produced them, each with a severity badge and a severity summary up top. Runs that opt in to **Save run documents** can also produce document-grade results — memos saved to the run's target knowledge base, listed under **Documents** on the receipt, where you can open them inline or download them. A run can also propose **memories** — durable notes it suggests keeping — listed under **Memories this run proposed**. Proposed memories can be kept or dismissed right on the receipt, or managed in the **Review** view. The Review view also lists **precedents** — recurring patterns the agent noticed — which you can dismiss or promote into a matter's context (you approve the final write). Results stream in live while the run works.

NOTIFICATIONS

When a run finishes, it posts to the **Notifications** view — title, summary, and a link straight to the session's receipt. Unread notifications carry a dot; **Mark read** clears them one at a time.

COST CAPS

Every way of starting a run — run now, schedules, watches — accepts an optional cost cap in USD. A run that reaches its cap stops, and the receipt records that as the terminal reason. Combined with receipts and notifications, you always know what ran, what it cost, and what it produced.

Tabular review

WHAT IT IS

Tabular review runs the same set of questions across many documents at once and returns the answers in a cited grid — one row per document, one column per question. Use it when you need to compare how different agreements, filings, or contracts answer the same question, or when you want to extract structured facts from a large document set without reading each file individually.

Each cell in the grid carries a confidence signal (**high**, **medium**, or **low**) and links to the source passage in the original document. You can export the finished grid to Excel or CSV for use outside Donna.

STARTING A NEW REVIEW

Open **Tabular** in the sidebar, then click **New review**. The builder page is split into two panels: **Documents** on the left and **Columns** on the right.

To add documents, switch between the **From a matter** and **Upload** tabs. When using a matter, pick the matter from the dropdown and check each file you want to include. When uploading, drag files onto the drop zone; they are ingested immediately and added to the selection as they finish processing. Selected documents appear as removable chips below the picker.

DEFINING COLUMNS

The **Columns** panel offers two modes, toggled with the **Define columns / Use a table skill** control:

- **Define columns** — build your own grid ad hoc. Each column has a **Column name** (the heading that appears in the grid) and a **Column question** (what Donna should extract, e.g. "Which state's law governs?"). Use the arrow buttons to reorder columns and the × button to remove one. Click **Add column** to add more. Column names must be unique.
- **Use a table skill** — pick a pre-built table skill from the list. The skill defines the columns; you cannot edit them here. The exact column count and estimated cost appear in the cost preview before you commit.

When defining columns manually, each column also has an optional **Min. model tier** setting (tiers 1–5, or none). Setting a minimum tier tells Donna to use at least that capability level when answering that column's question. Leave it at **None** to let the system choose automatically. Higher tiers may increase the estimated cost.

Each manually-defined column also has an **Ensemble verification** checkbox. When enabled, the answers in that column are independently cross-checked before being accepted, and when you open one of a verified cell's citations, the source view shows a green ✓ **Verified** chip. Ensemble-verified cells cost more; the cost preview shows the count of ensemble-verified cells and the premium as a separate line.

RUNNING THE REVIEW

The status bar at the bottom of the builder shows the current scope — for example **4 docs × 3 cols = 12 cells** — so you can gauge the size of the run before committing. When you are ready, click **Preview cost**.

A cost preview modal appears, showing the total cell count and an estimated cost in USD broken down by model tier. The estimate is approximate. Click **Run review** to start, or **Cancel** (or press Escape) to go back and adjust the setup.

Once started, the review runs in the background. The grid page opens automatically and shows a live progress indicator while cells are being processed. If a run takes longer than expected, a note appears suggesting you reload to check again. You can click **Cancel** on the running review page to stop it early.

READING THE GRID

When the review completes, the grid shows one row per document and one column per question. Each cell displays a coloured dot indicating confidence — green for high, amber for medium, red for low, and grey for a failed cell — followed by a short excerpt of the answer. A number on the right edge of a cell indicates how many source citations back the answer.

A summary line above the grid shows the total cell count and how many cells failed. If any cell shows **(failed)**, the cell detail panel (see below) will display the error reason when available.

NAVIGATING FROM A CELL TO ITS SOURCE

Click any cell to open the **cell detail panel** — a slide-in drawer on the right side of the screen. The panel shows the full answer text, the confidence level, and a list of citations. Each citation is a button labelled **Source · p.N**, with a short preview of the supporting passage below it.

Clicking a citation opens the **document panel** — the same viewer used elsewhere in Donna — which loads the source document and jumps to the cited page and passage, with a green ✓ **Verified** chip when the cell was ensemble-verified. Press Escape or click the × button to close the cell detail panel.

EXPORTING RESULTS

On any completed review, click the **Export ▾** button at the top right of the grid and choose **Excel (.xlsx)** or **CSV (.csv)**. The file downloads immediately.

HISTORY AND RESUMING PAST REVIEWS

All past and in-progress reviews are listed on the **Tabular** index page. Each row shows the run's status (completed, running, pending, failed, or cancelled), the number of documents and columns, an estimated cost if available, and the date the review was created. Click any row to open that review's grid — completed reviews show their full results immediately, while still-running reviews continue polling for updates in the background.

Research

SEARCHING CASE LAW

The **Research** tab in the sidebar opens a dedicated workspace for U.S. case law. Donna looks up court opinions through **CourtListener**, a free repository of millions of decisions. Search by topic (e.g. *"Chevron deference"*) or by case name (e.g. *"Brown v. Board of Education"*), and narrow results with the court and sort-order filters.

Click any result to open the full opinion in the document panel and read the decision. Once an opinion is open, use **Find in opinion** to jump to a phrase within it.

VERIFYING CITATIONS

The **Verify citations** tool takes text containing reporter citations (for example a paragraph from a brief) and checks each citation against CourtListener, linking the ones it finds to the matching case. It is a fast way to confirm that the authorities in a document are real and correctly cited.

RESEARCH FROM CHAT

You can also simply ask the **Assistant** to research for you — for example, *"Find a landmark Supreme Court case on free speech and cite it."* When Donna consults case law to answer, it shows a **"⚖ sources consulted"** note beneath the reply listing the cases it pulled in, each linking out to CourtListener so you can read the original.

Research in chat works best on a capable model — pick **smart** in the composer's model selector. Smaller local models can struggle to use the case-law tool reliably and may return an empty reply.

AVAILABILITY

Case-law research is enabled by your administrator with a CourtListener API token, so each operator brings their own key. If the Research tab says it isn't enabled, ask your administrator to add a token for this deployment.

Knowledge bases

WHAT IS A KNOWLEDGE BASE?

A **Knowledge base** is a named collection of documents that Donna can search when answering your questions. Instead of relying only on what the AI model already knows, Donna retrieves relevant passages from your uploaded documents and uses them to ground its answers — a technique called retrieval-augmented generation (RAG). This means answers can draw on your specific contracts, policies, research, or other materials rather than general training data.

CREATING A KNOWLEDGE BASE

Knowledge bases live inside **Projects**. Open a Project and find its **Knowledge** section: the picker there lets you link an existing KB or create a new one. When creating, you give the KB a name — pick something descriptive, such as "Acme Master Agreements" or "Company Policies Q4". Linking a KB to a Project makes its documents available to every chat in that Project.

Click a linked KB (or its **Manage** link) to open the knowledge base's own page. From there you can add files, review ingestion status, rename it, and tune its retrieval settings.

RENAMING AND DESCRIBING A KNOWLEDGE BASE

Open a KB and click **Rename** in the header. A dialog lets you update both the **Name** and an optional **Description**. Click **Save** to apply the changes. Press **Cancel** or hit Escape to close without saving.

UPLOADING DOCUMENTS

Inside a KB, the **Files** section shows all attached documents. If the KB is empty you will see a drop zone — drag files onto it or click to open a file picker. Once files are present, use the **+ Add file** button to upload more.

After upload, each file moves through ingestion states: **Pending** → **Processing** → **Ready**. Donna polls automatically and attaches the file once it is ready. If a file is still processing after several minutes, a **Refresh** prompt appears. Files that cannot be ingested show a **Failed** status with an error message.

Recommended format: PDF. PDF files ingest reliably. Some other file types — including plain `.txt` files — may be rejected with an "unsupported type" error. If a file fails to ingest, try converting it to PDF before uploading.

To remove a document, click the × icon next to the file row. This detaches the file from the KB; it will no longer be searched when Donna retrieves context. Ready files also show a **Download** link so you can retrieve the original document directly from the KB.

HOW RETRIEVAL WORKS

When you chat within a Project that has a linked knowledge base, Donna searches those documents for the passages most relevant to your query. Matching chunks are included in the context sent to the AI model, so the answer is grounded in your actual materials rather than model knowledge alone.

Each KB also exposes a **Hybrid alpha** control (in the **Advanced** section of the KB detail view). This slider blends two retrieval strategies: pure vector similarity search at one end, and full-text search (FTS) at the other. The default balance works well for most document types; you can tune it if you find one strategy returns better results for your content.

Answers that draw from a knowledge base cite their sources so you can verify exactly which document and passage the answer came from. See the **Trust & citations** page for a full explanation of how citations work in Donna.

ARCHIVING A KNOWLEDGE BASE

When a knowledge base is no longer needed, open it and click **Archive** in the header. Archived KBs are removed from the active list and will not be searched. This is useful for keeping the Knowledge view tidy without permanently deleting documents.

Models

CHOOSING A MODEL PER MESSAGE

Every message you send uses a model alias — a named routing slot that the backend resolves to a concrete model at request time. You choose your alias in the composer's model picker before sending. The picker shows only chat-capable aliases; embedding and other non-chat model types are filtered out automatically.

The two most common aliases you will see are:

- **smart** — routes to the highest-capability cloud model currently configured. Choose this when quality and depth matter more than speed.
- **fast** (or other lower-tier aliases) — routes to a faster, lighter model. Useful for quick drafts, summaries, or when you want lower latency.

The picker also surfaces any installed local models (served via Ollama). Local aliases are labelled **Local** in the picker. They run on-device and never send data to a cloud provider, making them a good choice for sensitive or offline work.

Some matters enforce a minimum inference tier. When you open a matter that requires a higher-tier model than your current selection, Donna automatically promotes your choice to a model that satisfies the floor — preferring **smart** when it qualifies, otherwise the highest-tier option available — so you never accidentally send a message with an underpowered model.

SETTINGS → MODELS

Workspace-level model routing is managed in **Settings → Models**, reachable from the sidebar. The page has three sections.

INFERENCE CATEGORIES

The **Inference categories** section lists every named alias in the system. Each row shows the category name, whether it is backed by a cloud or local model, and its current backing model label.

- **Admins** see a drop-down on every row and can reassign a category to any available cloud or local model. Changes take effect immediately — the row shows **Saving...** while the request is in flight and **Saved** on success.
- **Non-admins** see the current backing model for each category but cannot change it. A note at the bottom of the section explains that changing model routing requires an admin account.

INSTALLED LOCAL MODELS

The **Installed local models** card lists every Ollama model detected on your system, showing the model's label and its full identifier. If no local models are found, the card explains how to install Ollama and pull a model with `ollama pull <model>`. Adding a new model via Ollama and reloading the page is all that is needed for it to appear here and become selectable in the composer.

PROVIDER KEYS

The **Provider keys** card lists each model provider the deployment knows about (Anthropic, OpenAI, and any others) with its key status — ✓ **Configured**, with the key's source and last four characters, or no key yet. Keys are write-only: Donna never shows a stored key, only its last four characters.

- **Admins** paste a key into the masked input on a row and press **Add key** (or **Replace key**). Changes apply immediately — no restart needed.
- Keys set by the deployment's **environment** can be taken over by saving a runtime key on the same row. Runtime keys can be revoked with a two-step confirm; environment-managed keys cannot be revoked from the UI.
- **Non-admins** see a note that provider keys are managed by your administrator.

Tools & connections

Beyond answering from your documents and knowledge bases, Donna's assistant can use **tools** — external services that look things up or take an action — when they help answer your request. Donna keeps you in control of when and whether a tool runs.

MCP TOOLS

MCP (Model Context Protocol) servers are collections of external tools your administrator connects to the deployment — for example a documentation or reference lookup. Under **Settings** → **MCP** you can see which servers are connected and enable the individual tools you want the assistant to be able to use.

APPROVING A TOOL IN CHAT

When the assistant wants to use a tool, it pauses and asks first: a card appears in the chat naming the tool and what it will do, with **Approve** and **Deny** buttons. Approve lets it run that one time; Deny skips it and the assistant continues without it. Tools that could change something are flagged as destructive. Nothing external happens without your click.

CONNECTIONS

Some tool servers require you to sign in with your own account. Under **Settings** → **Connections** you connect each one once; afterwards the assistant can use it on your behalf. If you start a chat that needs a connection you haven't made yet, the assistant shows a **Connect** prompt inline so you can link it on the spot.

WHERE ANSWERS CAME FROM

When a tool pulls in outside material — such as case law from CourtListener — Donna shows a **"sources consulted"** note beneath the answer listing those external sources with links. This is separate from the green / amber / red **verified-quote citations** described on the **Trust & citations** page: those check a quote against a source, while "sources consulted" simply records which outside materials the assistant looked at.

Trust & citations

CITATION PILLS

When Donna's answer is backed by documents in your knowledge base, numbered **citation pills** appear inline in the response text — one pill per supporting passage. Hovering or focusing a pill opens a small popover that shows the exact quoted text from the source, the filename, and the page number when the backend provides it.

The popover header is colour-coded to signal how confidently the passage was matched:

- **Green — Verified.** The passage was found in the source as an exact or normalised match, or confirmed by ensemble verification (multiple judges agreed it is supported).
- **Amber — Caveats.** The source partially supports the claim, or the match was made by a single paraphrase judge rather than a direct text search.
- **Red — Unverified.** The citation could not be confirmed against the source document.

Clicking a pill (or pressing **Enter** or **Space** when focused) opens the **Document panel** on the right side of the screen so you can read the full surrounding context in the source.

RECEIPTS

Every chat page has a **Receipts** button in the top-right corner. Opening it slides in a panel that lists every backend event recorded during the conversation — from your initial message, through any knowledge-base retrieval and model inference steps, to the final answer.

Each row is labelled by kind: **message**, **retrieval**, **inference**, **skill**, **audit**, or **error**. Retrieval rows show how many chunks were pulled and from how many knowledge bases. Inference rows show the model name, provider, token counts, and latency. Selecting **Details** on any row expands the raw event payload for deeper inspection.

You can filter the list by kind using the chips at the top of the drawer, and export the full event log as a JSONL file with the **Export** link in the drawer header.

ANONYMIZATION INDICATOR

When a cloud model is used, your prompt passes through an **anonymization layer** before it leaves your environment — identifiers are stripped at the edge. Donna surfaces this as an **Anonymized** badge (with a shield-check icon) that appears beside the answer whenever the backend confirms that anonymization was applied to that inference call.

The same signal appears inside the Receipts drawer: each inference row shows either **Anonymized** (green) or **No anonymization** (grey) depending on what the backend recorded for that event. Local models, which run entirely on-device, do not send data outbound at all and therefore do not carry this indicator.

STANDING DISCLAIMER

Every chat page and the landing composer both display the notice: **"AI can make mistakes. Answers are not legal advice."** This reflects a core design principle of Donna: the system surfaces sources and verification signals so you can judge the answer yourself — it does not replace the professional judgement of a qualified lawyer. Always verify critical information against the underlying documents before acting on it.

Fiduciary receipts

Donna keeps an **honest provenance record on every answer**: for each reply, what sources it used, exactly what it quoted from each, whether that quote was actually found in the real source, and — for cases — signals about how later courts have treated it. You never have to take the answer on faith; you can trace any claim back to where it came from.

THE FOUR TRUST STATES

Every assistant answer carries a small **trust pill** that tells you, at a glance, how well its quoted claims held up when Donna checked them against their sources:

- **Fiduciary-grade (green).** The answer made sourced claims and every quoted claim was matched against its original source.
- **Supported (amber).** The claims are backed by the sources in substance — verified by meaning rather than an exact quote match.
- **Needs review (red).** At least one quoted claim could not be confirmed in its source. Read the flagged quote against the original before relying on it.
- **No sourced claims (neutral grey).** This answer did not quote or rely on a specific source. It is deliberately *never* shown as green: with nothing to verify, "verified" would be misleading — nothing to verify is not the same as verified.

Try each state below — pick what happened when Donna verified an answer, and see the exact pill you would get. Toggling "made sourced claims" off shows the honesty rule in action.

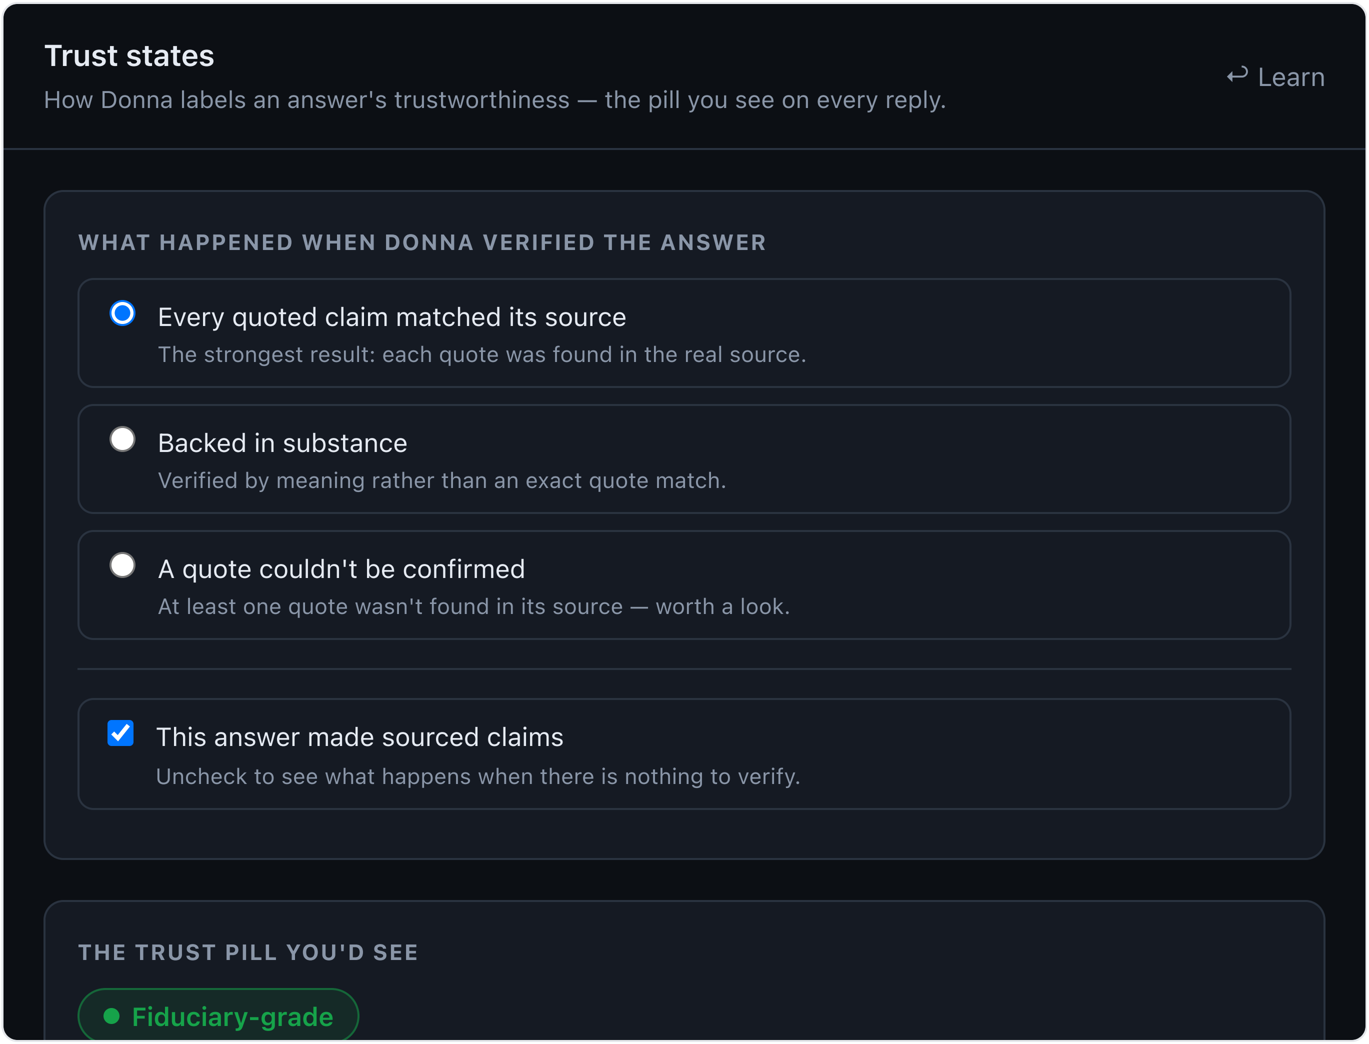

Open full-screen ↗

THE RECEIPT & CITATION LEDGER

Click the trust pill to open the **fiduciary receipt** — the citation ledger for that answer. It lists each source the assistant actually read, the exact passage it quoted, and a verification chip for that quote, so you can **trace any claim to its source** in one click. Knowledge-base sources open in the document panel; cases open the opinion; statutes and regulations link out to the authority. The ledger records *what was read and whether the quote matched*, referenced by identifier and passage location — not by stashing copies of the source text.

Beneath the answer-level trust pill, the engine buckets each individual citation into pass / supported / fail — a different layer from the four turn-level states above. Explore the citation ledger ↗ · the fiduciary gate ↗

CASE TREATMENT (VALIDITY)

Where an answer relies on case law, the receipt can show **derived treatment signals** — how often the case has been cited and whether later courts followed, distinguished, or criticised it — each with a trace to the citing case. This is **derived, not editorial**: it is a signal to guide your own reading, *not* an authoritative citator, and Donna never colours a case "good" or "bad" law for you.

Explore how treatment signals are derived ↗

MATTER SESSIONS — THE AUDIT TIMELINE

When Donna runs an autonomous matter session, its receipt is the same fiduciary record at the matter level — a timeline of who did what, on whose behalf, at what cost, alongside a session-wide trust pill answering the one question a reviewer cares about: **is this work product fiduciary-grade?**

Explore a governed matter session ↗

AUTHORITATIVE SOURCES

The Research page lists the primary-law sources this instance can reach — U.S. case law (CourtListener), the U.S. Code and CFR (GovInfo), public filings (SEC EDGAR), and EU legislation and CJEU case law (EUR-Lex). Each is retrieved *and* character-verified through the same governed path. Availability depends on how your operator has configured the instance, and EU law is currently fetch-by-identifier (CELEX) only; a source that isn't configured shows as unavailable rather than silently disappearing.

Explore the authoritative-source registry ↗

TAKE IT WITH YOU — PROVENANCE EXPORT

From any chat turn or matter session you can **export the provenance record** — a structured JSON file and a printable Markdown copy of the whole sourcing trail. It is labelled honestly for what it is: a faithful copy of the record, **not a cryptographically signed attestation**. It is designed so that, when a signed export becomes available, the same button can produce it.

UNDER THE HOOD

Want to see how a quote actually gets verified? The LQ-AI engine runs a four-stage character-fidelity cascade on every quoted passage. Drill into how a quote is verified ↗

The interactive walkthroughs linked above are **illustrative** — self-contained examples of how the LQ-AI engine works, not your live data — and each links onward to the real backend module or decision record it depicts.